

#Importing Liraries

Importing Libraries

```
import warnings
warnings.filterwarnings("ignore")
```

```
from tqdm import tqdm
import shutil
import os
import numpy as np
import pandas as pd
from datetime import datetime
import matplotlib
matplotlib.use('nbAgg')
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import pickle
import random
import joblib
from scipy.stats import randint as sp_randint
from scipy.stats import uniform
from scipy.sparse import hstack
from wordcloud import WordCloud
```

Utilities

```
#from viz_utils import *
#from custom_transformers import *
#from ml_utils import *
```

DataPrep

```
import re
from nltk.corpus import stopwords
from nltk.stem import RSLPStemmer
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.preprocessing import Normalizer
```

Modeling

```
from sklearn.model_selection import train_test_split
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.calibration import CalibratedClassifierCV
from sklearn.linear_model import SGDClassifier, LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, StackingClassifier, VotingClassifier
```

#Metrics

```
from sklearn.metrics import log_loss, accuracy_score, confusion_matrix, f1_score
```

```
#Importing the Libraries
```

```
from tensorflow.keras.layers import LeakyReLU
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.models import load_model
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
cd /content/drive/MyDrive/case study
```

```
/content/drive/MyDrive/case study
```

```
X_train = pickle.load(open('X_train.pkl', 'rb'))
X_test = pickle.load(open('X_test.pkl', 'rb'))
```

```
X_train.head()
```

	order_status	payment_sequential	payment_type	payment_installments	payment_value	customer_zip_code_prefix
28832	delivered	1	boleto	1	95.21	13211
28103	delivered	1	credit_card	1	54.28	8215
76931	delivered	1	credit_card	1	171.62	22745
17603	delivered	1	credit_card	10	139.55	35519
8208	delivered	1	boleto	1	82.91	20261

```
X_tr = pickle.load(open('X_tr.pkl', 'rb'))
X_te = pickle.load(open('X_te.pkl', 'rb'))
y_train = pickle.load(open('y_train.pkl', 'rb'))
y_test = pickle.load(open('y_test.pkl', 'rb'))
```

```
# load the model from file
encoder = load_model('encoder.h5')
```

```
# encode the train data
X_train_encode = encoder.predict(X_tr)
```

```

print('X_train_encode',X_train_encode.shape)
# encode the test data
X_test_encode = encoder.predict(X_te)
print('X_test_encode',X_test_encode.shape)

```

WARNING:tensorflow:No training configuration found in the save file, so the model was *not* compiled. Compile it manually.

X_train_encode (77804, 112)

X_test_encode (19451, 112)

```

def confusion_matrices_plot(y_real, y_pred, y_test,y_test_pred,name):
    # representing confusion matrix in heatmap format
    # https://seaborn.pydata.org/generated/seaborn.heatmap.html
    cmap_ = sns.light_palette("#425a90",as_cmap=True)
    cmap = sns.light_palette("#000000",as_cmap=True)
    C1 = confusion_matrix(y_real,y_pred)
    C2 = confusion_matrix(y_test,y_test_pred)

    fig,ax = plt.subplots(1, 2, figsize=(15,5))
    ax1 = sns.heatmap(C1, annot=True, cmap=cmap, fmt=".2f", ax = ax[0])
    ax1.set_xlabel('Predicted Class')
    ax1.set_ylabel('Original Class')
    ax1.set_title("Train Confusion matrix")

    ax2 = sns.heatmap(C2, annot=True, cmap=cmap_, fmt=".2f", ax = ax[1])
    ax2.set_xlabel('Predicted Class')
    ax2.set_ylabel('Original Class')
    ax2.set_title("Test Confusion matrix")
    plt.savefig(name, dpi=480, bbox_inches='tight')
    plt.show()

```

1.Linear SVM Model

```

# Training Logistic regression model and chekcing f1 score metric
alpha = [10 ** x for x in range(-5, 4)]
train_scores = [] # store train scores
test_scores = [] # store test scores

for i in alpha:
    lr = SGDClassifier(loss='hinge', penalty='l2', alpha=i, n_jobs=-1, random_state=25)
    lr.fit(X_train_encode,y_train)
    train_sc = f1_score(y_train,lr.predict(X_train_encode),average='macro')
    test_sc = f1_score(y_test,lr.predict(X_test_encode),average='macro')
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('Alpha = ',i,'Train Score',train_sc,'test Score',test_sc)


# plotting the scores vs parameters
plt.plot(np.log(alpha),train_scores,label='Train Score',color='black')

```

```
plt.plot(np.log(alpha), test_scores, label='Test Score', color='#425a90', alpha=0.7)
plt.xlabel('Alpha')
plt.ylabel('Score')
plt.grid()
plt.title('Alpha vs Score')
```

```
Alpha = 1e-05 Train Score 0.7426881274403494 test Score 0.711429758484188
Alpha = 0.0001 Train Score 0.7969280203375975 test Score 0.7590218036650437
Alpha = 0.001 Train Score 0.7779664010455214 test Score 0.7458017969782127
Alpha = 0.01 Train Score 0.8120023685210596 test Score 0.7670568438912866
Alpha = 0.1 Train Score 0.8454147613073184 test Score 0.7944185654229546
Alpha = 1 Train Score 0.8442628803938433 test Score 0.7930241824437074
Alpha = 10 Train Score 0.8218645706991656 test Score 0.7804407657723436
Alpha = 100 Train Score 0.4684095647897985 test Score 0.4712500475569102
Alpha = 1000 Train Score 0.4549814497049617 test Score 0.4565774667184236
```

```
Text(0.5, 1.0, 'Alpha vs Score')
```



```
# Parameter tuning of Logistic regression using RandomisedSearch CV technique
```

```
sgd = SGDClassifier(loss='hinge', n_jobs=-1, random_state=25)
```

```
prams={ 'alpha': [10 ** x for x in range(-5, 4)] }
```

```
random_cfl1 = RandomizedSearchCV(sgd,param_distributions=prams,verbose=10,scoring='f1_m
                                return_train_score=True)
```

```
random_cfl1.fit(X_train_encode,y_train)
```

```
print('mean test scores',random_cfl1.cv_results_['mean_test_score'])
```

```
print('mean train scores',random_cfl1.cv_results_['mean_train_score'])
```

Fitting 5 folds for each of 9 candidates, totalling 45 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 1 tasks | elapsed: 12.4s
```

```
[Parallel(n_jobs=-1)]: Done 4 tasks | elapsed: 28.5s
```

```
[Parallel(n_jobs=-1)]: Done 9 tasks | elapsed: 1.0min
```

```
[Parallel(n_jobs=-1)]: Done 14 tasks      | elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 21 tasks      | elapsed: 1.5min
[Parallel(n_jobs=-1)]: Done 28 tasks      | elapsed: 1.5min
[Parallel(n_jobs=-1)]: Done 37 tasks      | elapsed: 1.5min
[Parallel(n_jobs=-1)]: Done 45 out of 45 | elapsed: 1.6min finished
```

```
mean test scores [0.81283792 0.81544242 0.79118321 0.81786007 0.84249014 0.85295023
0.83787149 0.47277831 0.46096543]
```

```
mean train scores [0.81428754 0.81534436 0.79081021 0.81915052 0.84231144 0.85361772
0.8377479 0.47326957 0.46096551]
```

```
# printing best parameters and score
print("Best Parameters: ",random_cfl1.best_params_)
print("Best Score: ",random_cfl1.best_score_)
```

```
Best Parameters:  {'alpha': 1}
```

```
Best Score:  0.852950225246062
```

```
# Fitting LogisticRegression model on best parameters
sgd = SGDClassifier(loss='hinge', alpha=1, n_jobs=-1, random_state=25)
sgd.fit(X_train_encode,y_train)
```


```
y_train_pred = sgd.predict(X_train_encode)
y_test_pred = sgd.predict(X_test_encode)
```

```
# printing train and test scores
print('Train f1 score: ',f1_score(y_train,y_train_pred,average='macro'))
print('Test f1 score: ',f1_score(y_test,y_test_pred,average='macro'))
```

```
Train f1 score:  0.8442628803938433
```

```
Test f1 score:  0.7930241824437074
```

```
confusion_matrices_plot(y_train,y_train_pred,y_test,y_test_pred, 'svm.png')
```



```
#saving the model
filename = '/content/drive/MyDrive/case study/models/svm.sav'
joblib.dump(sgd, filename)
```


```
['/content/drive/MyDrive/case study/models/svm.sav']
```

2.Decision Tree

```
# Checking the variation of score with depth parameters of Decision Tree
depth = [3,10,50,100,250,500]
train_scores = []
test_scores = []
for i in depth:
    clf = DecisionTreeClassifier(max_depth=i,random_state=25)
    clf.fit(X_train_encode,y_train)
    train_sc = f1_score(y_train,clf.predict(X_train_encode),average='macro')
    test_sc = f1_score(y_test,clf.predict(X_test_encode),average='macro')
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('Depth = ',i,'Train Score',train_sc,'test Score',test_sc)

# plotting the score vs depth
plt.plot(depth,train_scores,label='Train Score',color='black')
plt.plot(depth,test_scores,label='Test Score',color='#425a90',alpha=0.6)
plt.xlabel('Depth')
plt.ylabel('Score')
plt.title('Depth vs Score')
plt.grid()
plt.show()
```

```
Depth = 3 Train Score 0.8410443354270485 test Score 0.7823336186969845
Depth = 10 Train Score 0.8944223553640378 test Score 0.8011274038641633
Depth = 50 Train Score 0.9607451887375799 test Score 0.7674886074762987
Depth = 100 Train Score 0.9607703670756125 test Score 0.7669917560036371
Depth = 250 Train Score 0.9607703670756125 test Score 0.7669917560036371
Depth = 500 Train Score 0.9607703670756125 test Score 0.7669917560036371
```



```

# Parameter tuning of DecisionTreeClassifier using RandomisedSearch CV technique
# https://medium.com/@mohtedibf/indepth-parameter-tuning-for-decision-tree-6753118a03c3
dt = DecisionTreeClassifier(random_state=25)

params = { "max_depth": sp_randint(3,500), "min_samples_split": sp_randint(50,200), "mi

random_cfl1 = RandomizedSearchCV(dt, param_distributions=params,verbose=10,scoring='f1_
                                return_train_score=True)
random_cfl1.fit(X_train_encode,y_train)

print('mean test scores',random_cfl1.cv_results_['mean_test_score'])
print('mean train scores',random_cfl1.cv_results_['mean_train_score'])

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done   1 tasks      | elapsed:   21.5s
[Parallel(n_jobs=-1)]: Done   4 tasks      | elapsed:   44.4s
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:   1.8min
[Parallel(n_jobs=-1)]: Done  14 tasks      | elapsed:   2.7min
[Parallel(n_jobs=-1)]: Done  21 tasks      | elapsed:   4.3min
[Parallel(n_jobs=-1)]: Done  28 tasks      | elapsed:   5.7min
[Parallel(n_jobs=-1)]: Done  37 tasks      | elapsed:   7.7min
[Parallel(n_jobs=-1)]: Done  46 tasks      | elapsed:   9.3min
[Parallel(n_jobs=-1)]: Done  50 out of   50 | elapsed: 10.0min finished

mean test scores [0.86400205 0.86391897 0.85812082 0.86370056 0.85713902 0.85693981
 0.84822633 0.86139325 0.86204052 0.85989181]
mean train scores [0.8758983  0.87667336 0.8817508  0.87602165 0.88223669 0.88235182
 0.89540709 0.88088749 0.87740286 0.87909035]

```

```

# printing best parameters and scores
print("Best Parameters: ",random_cfl1.best_params_)
print("Best Score: ",random_cfl1.best_score_)

```

```

Best Parameters:  {'max_depth': 135, 'min_samples_leaf': 28, 'min_samples_split': 193}
Best Score:  0.8640020459676931

```

```

# Fitting the model on best parameters
dt = DecisionTreeClassifier(max_depth = 135, min_samples_leaf = 28, min_samples_split =
dt.fit(X_train_encode,y_train)

y_train_pred = dt.predict(X_train_encode)
y_test_pred = dt.predict(X_test_encode)


# printing train test score
print('Train f1 score',f1_score(y_train,y_train_pred,average='macro'))
print('Test f1 score',f1_score(y_test,y_test_pred,average='macro'))

```

Train f1 score 0.8774934888896901

Test f1 score 0.8028745314241124

```
confusion_matrices_plot(y_train,y_train_pred,y_test,y_test_pred, 'dt.png')
```



```
#saving the model
```

```
filename = '/content/drive/MyDrive/case study/models/dt.sav'
```


```
joblib.dump(dt, filename)
```

```
['/content/drive/MyDrive/case study/models/dt.sav']
```

3. RandomForest Classifier

```
estimators = [5,10,50,100,250,500]
train_scores = []
test_scores = []
for i in estimators:
    clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=5, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=52, min_samples_split=120,
                                min_weight_fraction_leaf=0.0, n_estimators=i, n_jobs=-1, random_state=25, verbose=0)
    clf.fit(X_train_encode,y_train)
    train_sc = f1_score(y_train,clf.predict(X_train_encode),average='macro')
    test_sc = f1_score(y_test,clf.predict(X_test_encode),average='macro')
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('Estimators = ',i,'Train Score',train_sc,'test Score',test_sc)
plt.plot(estimators,train_scores,label='Train Score',color='black')
plt.plot(estimators,test_scores,label='Test Score',color='#425a90',alpha=0.7)
plt.xlabel('Estimators')
plt.ylabel('Score')
plt.title('Estimators vs score at depth of 5')
plt.grid()
```


Estimators = 5 Train Score 0.8482808531495858 test Score 0.7846015796896826
 Estimators = 10 Train Score 0.856353143130909 test Score 0.7895872404784336
 Estimators = 50 Train Score 0.8616414726517485 test Score 0.7945000965035711
 Estimators = 100 Train Score 0.8621485418017834 test Score 0.7962166986862518
 Estimators = 250 Train Score 0.8617113177255411 test Score 0.795732073874307
 Estimators = 500 Train Score 0.8619575273508977 test Score 0.7967678711988769

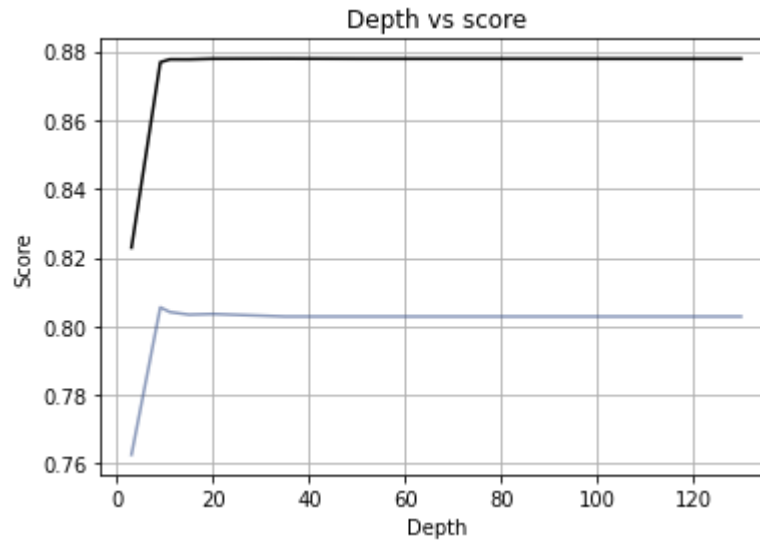


```

depths = [3,9,11,15,20,35,50,70,130]
train_scores = []
test_scores = []
for i in depths:
    clf = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=i, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=52, min_samples_split=120,
                                min_weight_fraction_leaf=0.0, n_estimators=115, n_jobs=-1, random_state=25,
                                verbose=0)
    clf.fit(X_train_encode, y_train)
    train_sc = f1_score(y_train, clf.predict(X_train_encode), average='macro')
    test_sc = f1_score(y_test, clf.predict(X_test_encode), average='macro')
    test_scores.append(test_sc)
    train_scores.append(train_sc)
    print('depth = ', i, 'Train Score', train_sc, 'test Score', test_sc)
plt.plot(depths, train_scores, label='Train Score', color='black')
plt.plot(depths, test_scores, label='Test Score', color='#425a90', alpha=0.6)
plt.xlabel('Depth')
plt.ylabel('Score')
plt.title('Depth vs score')
plt.grid()
plt.show()
  
```

depth = 3 Train Score 0.8229182137899087 test Score 0.7626631883898214
 depth = 9 Train Score 0.8768412814477038 test Score 0.8055227907760845
 depth = 11 Train Score 0.8776910412757344 test Score 0.8041546148784972
 depth = 15 Train Score 0.8776866740714329 test Score 0.8034038211809486
 depth = 20 Train Score 0.8778826837987468 test Score 0.803591094147101

depth = 35 Train Score 0.8779081890159492 test Score 0.8028553596843356
 depth = 50 Train Score 0.8778663113061898 test Score 0.8028553596843356
 depth = 70 Train Score 0.8778663113061898 test Score 0.8028553596843356
 depth = 130 Train Score 0.8778663113061898 test Score 0.8028553596843356



```
from sklearn.metrics import f1_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import f1_score
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint
from scipy.stats import uniform

param_dist = {"n_estimators":sp_randint(105,125),
              "max_depth": sp_randint(10,15),
              "min_samples_split": sp_randint(110,190),
              "min_samples_leaf": sp_randint(25,65)}

clf = RandomForestClassifier(random_state=25,n_jobs=-1)

rf_random = RandomizedSearchCV(clf, param_distributions=param_dist,
                               n_iter=5,cv=10,scoring='f1_macro',random_state=25,refit=True)

rf_random.fit(X_train_encode,y_train)
print('mean test scores',rf_random.cv_results_['mean_test_score'])
print('mean train scores',rf_random.cv_results_['mean_train_score'])
```

mean test scores [0.869705 0.86978126 0.86934033 0.86987059 0.86907663]
 mean train scores [0.87793813 0.87823902 0.87626858 0.87692678 0.87985437]

```
# printing best parameters and score
print("Best Parameters: ",rf_random.best_params_)
print("Best Score: ",rf_random.best_score_)
```

Best Parameters: {'max_depth': 13, 'min_samples_leaf': 49, 'min_samples_split': 165, 'n_estimators': 108}
 Best Score: 0.8698705892967556

```
# Fitting the model on best parameters
```

```
rf_classifier = RandomForestClassifier(max_depth = 13, min_samples_leaf = 49, min_samples_split = 10,  
                                     n_jobs=-1)  
rf_classifier.fit(X_train_encode,y_train)
```

```
y_train_pred = rf_classifier.predict(X_train_encode)
```

```
y_test_pred = rf_classifier.predict(X_test_encode)
```

```
# printing train and test scores
```

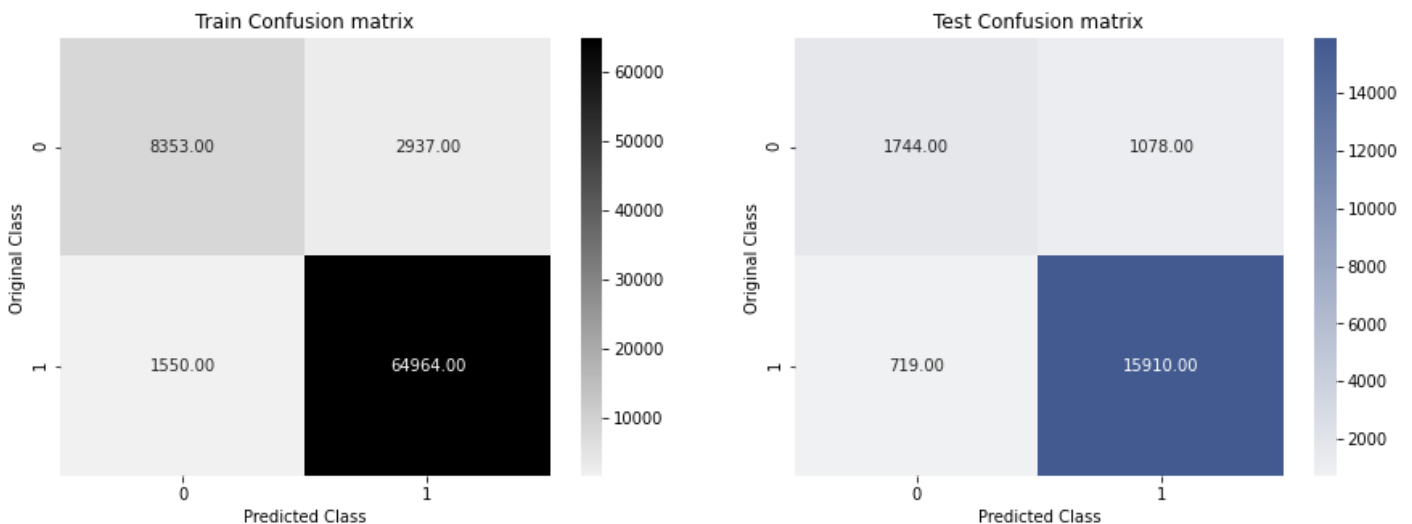
```
print('Train f1 score',f1_score(y_train,y_train_pred,average='macro'))
```

```
print('Test f1 score',f1_score(y_test,y_test_pred,average='macro'))
```

Train f1 score 0.8774487288720161

Test f1 score 0.8032629907003228

```
confusion_matrices_plot(y_train,y_train_pred,y_test,y_test_pred,'rf.png')
```



```
# save the model to disk
```

```
filename = '/content/drive/MyDrive/case study/models/randomf.sav'
```

```
joblib.dump(rf_classifier, filename)
```

```
['/content/drive/MyDrive/case study/models/randomf.sav']
```

4.LGBM

```
# Variation of score with estimators used in LGBM with other parameters set to default
```

```
estimators = [1,3,5,10,50,100,250,500,1000]
```

```
train_scores = []
```

```
test_scores = []
```

```
for i in estimators:
```

```
    clf = LGBMClassifier(n_estimators=i, n_jobs=-1,random_state=25)
```

```
    clf.fit(X_train_encode,y_train)
```

```

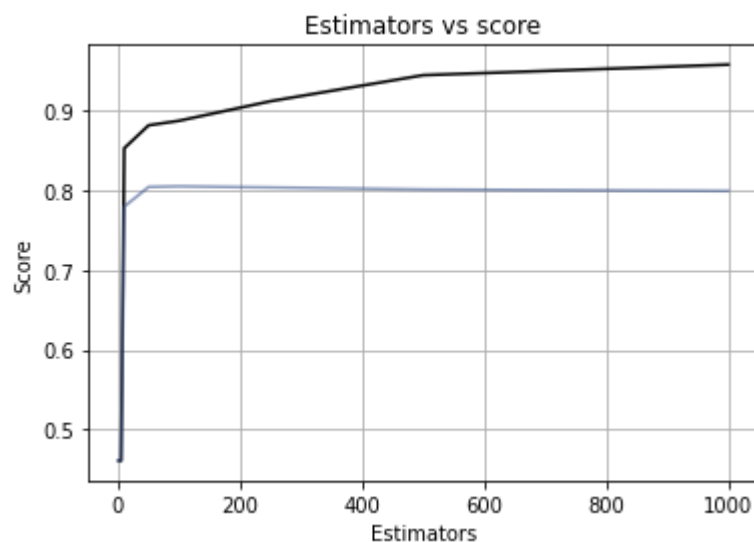
train_sc = f1_score(y_train,clf.predict(X_train_encode),average='macro')
test_sc = f1_score(y_test,clf.predict(X_test_encode),average='macro')
test_scores.append(test_sc)
train_scores.append(train_sc)
print('Estimators = ',i,'Train Score',train_sc,'test Score',test_sc)
plt.plot(estimators,train_scores,label='Train Score',color='black')
plt.plot(estimators,test_scores,label='Test Score',color='#425a90',alpha=0.6)
plt.xlabel('Estimators')
plt.ylabel('Score')
plt.title('Estimators vs score')
plt.grid()

```

```

Estimators = 1 Train Score 0.4608849900913261 test Score 0.46089246119733923
Estimators = 3 Train Score 0.4608849900913261 test Score 0.46089246119733923
Estimators = 5 Train Score 0.4608849900913261 test Score 0.46089246119733923
Estimators = 10 Train Score 0.8527908646965197 test Score 0.7795356610146007
Estimators = 50 Train Score 0.8815653668527856 test Score 0.8042926427631747
Estimators = 100 Train Score 0.887269982513814 test Score 0.8050377016191237
Estimators = 250 Train Score 0.911807695468607 test Score 0.8036228169737664
Estimators = 500 Train Score 0.9444291804782906 test Score 0.8010753511424802
Estimators = 1000 Train Score 0.9576994317083201 test Score 0.7991633195885934

```



[https://www.analyticsvidhya.com/blog/2016/03/complete-guide-parameter-tuning-xgboost-](https://www.analyticsvidhya.com/blog/2016/03/complete-guide-parameter-tuning-xgboost/)
Parameter tuning of the LGBM parameters using RandomSearch CV

```
x_cfl=LGBMClassifier(random_state=25,n_jobs=-1)
```

```

prams={
    'learning_rate':[0.001,0.01,0.03,0.05,0.1,0.15,0.2],
    'n_estimators':[1,3,5,10,50,100,250,500,1000],
    'max_depth':[3,5,10,15,20,50],
    'colsample_bytree':[0.1,0.3,0.5,1],
    'subsample':[0.1,0.3,0.5,1]
}

```

```

random_cfl1=RandomizedSearchCV(x_cfl,param_distributions=prams,verbose=10,n_jobs=-1,random_state=25,
                               return_train_score=True)
random_cfl1.fit(X_train_encode,y_train)

```

```
print('mean test scores',random_cfl1.cv_results_['mean_test_score'])
print('mean train scores',random_cfl1.cv_results_['mean_train_score'])
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done   1 tasks      | elapsed:   1.1min
```

```
[Parallel(n_jobs=-1)]: Done   4 tasks      | elapsed:   2.2min
```

```
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:   3.1min
```

```
[Parallel(n_jobs=-1)]: Done  14 tasks      | elapsed:   3.5min
```

```
[Parallel(n_jobs=-1)]: Done  21 tasks      | elapsed:   3.7min
```

```
[Parallel(n_jobs=-1)]: Done  28 tasks      | elapsed:   3.8min
```

```
[Parallel(n_jobs=-1)]: Done  37 tasks      | elapsed:   4.1min
```

```
[Parallel(n_jobs=-1)]: Done  46 tasks      | elapsed:   4.3min
```

```
[Parallel(n_jobs=-1)]: Done  50 out of  50 | elapsed:   4.4min finished
```

```
mean test scores [0.86547853 0.87051832 0.87032389 0.86893712 0.46088499 0.86847912
 0.86961699 0.46088499 0.8152882  0.46088499]
```

```
mean train scores [0.93449821 0.88242659 0.87972026 0.88836053 0.46088499 0.89321307
 0.87461157 0.46088499 0.82309595 0.46088499]
```

```
# printing best parameters and score
```

```
print("Best Parameters: ",random_cfl1.best_params_)
```

```
print("Best Score: ",random_cfl1.best_score_)
```

```
Best Parameters:  {'subsample': 0.5, 'n_estimators': 500, 'max_depth': 10,
'learning_rate': 0.01, 'colsample_bytree': 0.1}
```

```
Best Score:  0.870518321178464
```

```
# Fitting the model on best parameters
```

```
lgbm = LGBMClassifier(n_estimators=500, max_depth=10,subsample=0.5,learning_rate=0.01,c
                      n_jobs=-1)
```

```
lgbm.fit(X_train_encode,y_train)
```

```
y_train_pred = lgbm.predict(X_train_encode)
```

```
y_test_pred = lgbm.predict(X_test_encode)
```

```
# printing train and test scores
```

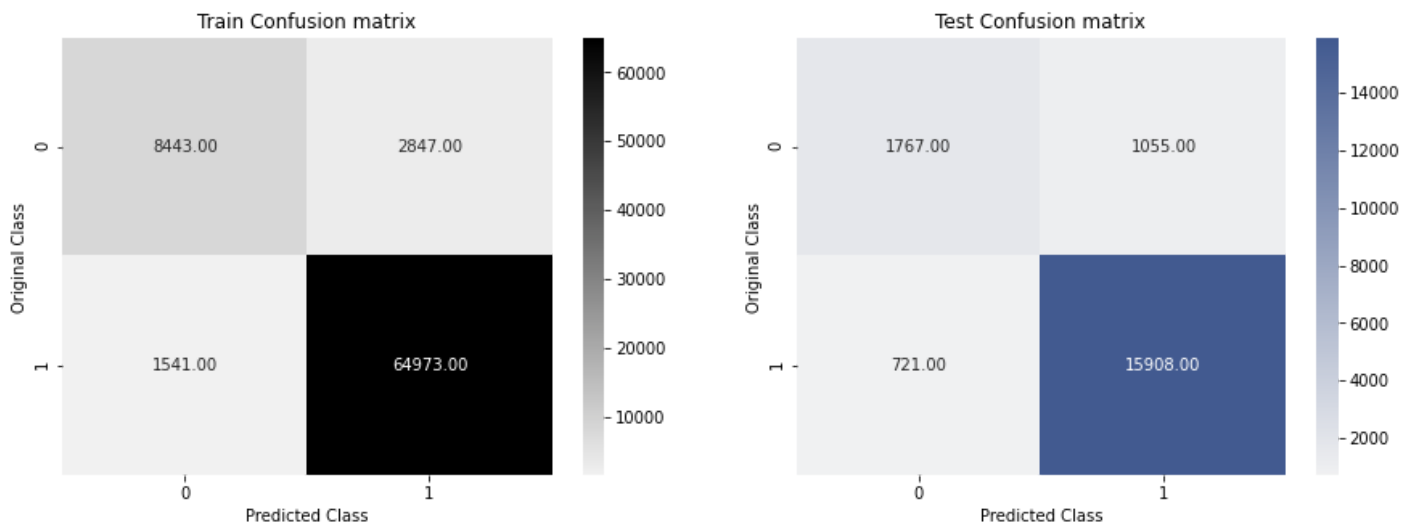
```
print('Train f1 score',f1_score(y_train,y_train_pred,average='macro'))
```

```
print('Test f1 score',f1_score(y_test,y_test_pred,average='macro'))
```

```
Train f1 score 0.8805369929197495
```

```
Test f1 score 0.8063334961377676
```

```
confusion_matrices_plot(y_train,y_train_pred,y_test,y_test_pred, 'lgbm.png')
```



```
filename = '/content/drive/MyDrive/case study/models/lgbm.sav'
joblib.dump(lgbm, filename)
```

```
['/content/drive/MyDrive/case study/models/lgbm.sav']
```

5.XGB Classifier

```
from xgboost import XGBClassifier
import random
clf = XGBClassifier()
param_dist = {"n_estimators":sp_randint(105,125),
              "eta":[round(random.uniform(0.1, 1.0 ), 1) for i in range(0,20)]
              }
xg_random = RandomizedSearchCV(clf, param_distributions=param_dist,
                               n_iter=5,cv=3,scoring='f1_macro',random_state=25,ret

xg_random.fit(X_train_encode,y_train)
print('mean test scores',xg_random.cv_results_['mean_test_score'])
print('mean train scores',xg_random.cv_results_['mean_train_score'])
```

```
mean test scores [0.86772465 0.86785741 0.86768704 0.8673134 0.86747848]
mean train scores [0.87655227 0.87655871 0.8760196 0.87594627 0.87599905]
```

```
print(xg_random.best_estimator_)
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, eta=0.3, gamma=0,
              learning_rate=0.1, max_delta_step=0, max_depth=3,
              min_child_weight=1, missing=None, n_estimators=123, n_jobs=1,
              nthread=None, objective='binary:logistic', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=None, subsample=1, verbosity=1)
```

```
clf = XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=1, eta=0.3, gamma=0,
                    learning_rate=0.1, max_delta_step=0, max_depth=3,
                    min_child_weight=1, missing=None, n_estimators=123, n_jobs=1,
                    nthread=None, objective='binary:logistic', random_state=0,
                    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                    silent=None, subsample=1, verbosity=1)
```


```
clf.fit(X_tr,y_train)
y_train_pred = clf.predict(X_tr)
y_test_pred = clf.predict(X_te)
```

```
# printing train and test scores
print('Train f1 score',f1_score(y_train,y_train_pred,average='macro'))
print('Test f1 score',f1_score(y_test,y_test_pred,average='macro'))
```

Train f1 score 0.8062058946636459

Test f1 score 0.7818943463057906

```
confusion_matrices_plot(y_train,y_train_pred,y_test,y_test_pred, 'xgboost.png')
```



```
filename = '/content/drive/MyDrive/case study/models/xgboost.sav'
joblib.dump(clf, filename)
```

```
['/content/drive/MyDrive/case study/models/xgboost.sav']
```

6. Adaboost Model

```
# https://www.analyticsvidhya.com/blog/2016/03/complete-guide-parameter-tuning-xgboost-
# Parameter tuning of the LGBM parameters using RandomSearch CV
x_cfl=AdaBoostClassifier(random_state=25)

prams={
```

```

    'learning_rate':[0.001,0.01,0.03,0.05,0.1,0.15,0.2],
    'n_estimators':[1,3,5,10,50,100,250,500,1000]
}
random_cfl1=RandomizedSearchCV(x_cfl1,param_distributions=params,verbose=10,n_jobs=-1,random_state=42,
                              return_train_score=True)
random_cfl1.fit(X_train_encode,y_train)

print('mean test scores',random_cfl1.cv_results_['mean_test_score'])
print('mean train scores',random_cfl1.cv_results_['mean_train_score'])

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done   1 tasks      | elapsed:    7.0s
[Parallel(n_jobs=-1)]: Done   4 tasks      | elapsed:   13.3s
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:   19.8s
[Parallel(n_jobs=-1)]: Done  14 tasks      | elapsed:  2.4min
[Parallel(n_jobs=-1)]: Done  21 tasks      | elapsed: 23.9min
[Parallel(n_jobs=-1)]: Done  28 tasks      | elapsed: 24.3min
[Parallel(n_jobs=-1)]: Done  37 tasks      | elapsed: 36.6min
[Parallel(n_jobs=-1)]: Done  46 tasks      | elapsed: 49.5min
[Parallel(n_jobs=-1)]: Done  50 out of  50 | elapsed: 68.2min finished

mean test scores [0.81617065 0.81548033 0.81406621 0.81406621 0.81073878 0.81559259
 0.81273864 0.85922691 0.81544965 0.86096781]
mean train scores [0.81663367 0.81649262 0.81518355 0.81518355 0.8111668  0.81613359
 0.812844  0.86040939 0.81619281 0.86273996]

```

```

# printing best parameters and score
print("Best Parameters: ",random_cfl1.best_params_)
print("Best Score: ",random_cfl1.best_score_)

```

```

Best Parameters:  {'n_estimators': 500, 'learning_rate': 0.05}
Best Score:  0.8609678086080766

```

```

# Fitting the model on best parameters
ada = AdaBoostClassifier(n_estimators=500, learning_rate=0.05, random_state=25)
ada.fit(X_train_encode,y_train)

y_train_pred = ada.predict(X_train_encode)
y_test_pred = ada.predict(X_test_encode)

# printing train and test scores
print('Train f1 score',f1_score(y_train,y_train_pred,average='macro'))
print('Test f1 score',f1_score(y_test,y_test_pred,average='macro'))

```

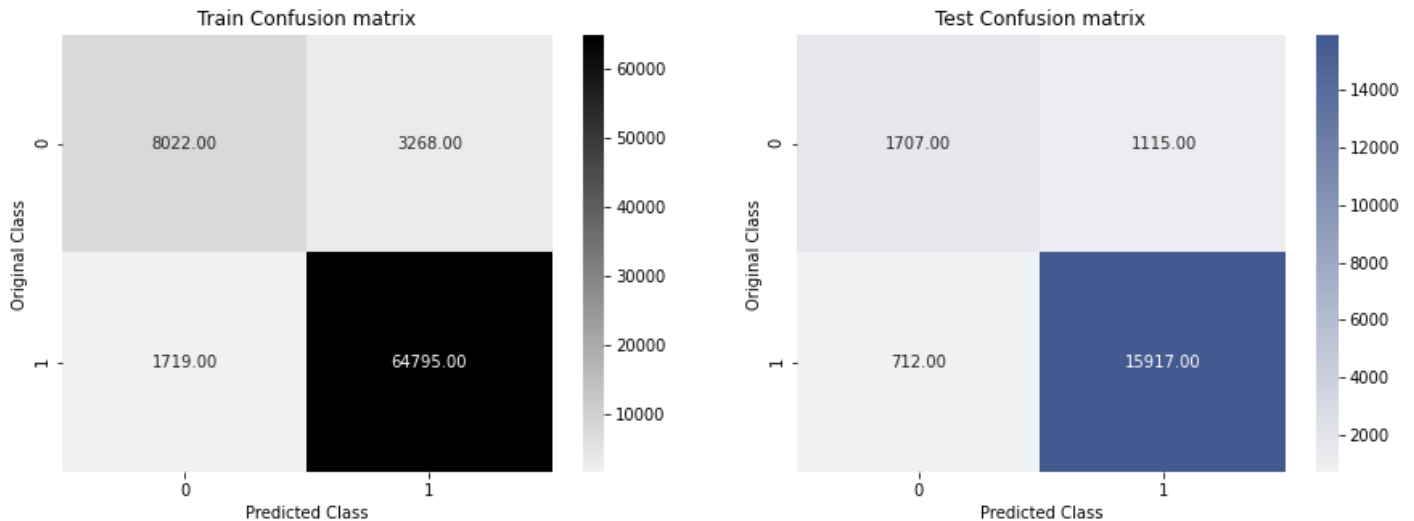
```

Train f1 score 0.8629085003364181
Test f1 score 0.7985629708732479

```



```
confusion_matrices_plot(y_train,y_train_pred,y_test,y_test_pred,'ada.png')
```



```
filename = '/content/drive/MyDrive/case study/models/ada.sav'
joblib.dump(ada, filename)
```

```
['/content/drive/MyDrive/case study/models/ada.sav']
```

7.Catboost Classifier

```
from catboost import CatBoostClassifier
learning_rate = [0.001,0.01,0.1]
max_depth = [5,10,15]

#Loop for the Tuning
for lr in learning_rate:
    for n in max_depth:
        clf = CatBoostClassifier(learning_rate=lr,max_depth = n,verbose = False)

        clf.fit(X_train_encode,y_train)

        #y_pred = (clf.predict_proba(X_test)[: , 1] >= 0.21).astype('int')
        Y_pred = clf.predict(X_train_encode)
        y_pred = clf.predict(X_test_encode)
        print(f"Train f1 score at learning rate {lr} and max_depth {n} is {f1_score(y_t
        print(f"Test f1 score at learning rate {lr} and max_depth {n} is {f1_score(y_te
        print("-"*50)
```

Train f1 score at learning rate 0.001 and max_depth 5 is 0.8671317844968727

Test f1 score at learning rate 0.001 and max_depth 5 is 0.8037009139263656

Train f1 score at learning rate 0.001 and max_depth 10 is 0.873491923891086

Test f1 score at learning rate 0.001 and max_depth 10 is 0.8052594369922508

Train f1 score at learning rate 0.001 and max_depth 15 is 0.8843754862374552

Test f1 score at learning rate 0.001 and max_depth 15 is 0.8045287713262363

Train f1 score at learning rate 0.01 and max_depth 5 is 0.8763461174771061

Test f1 score at learning rate 0.01 and max_depth 5 is 0.805690281820876

Train f1 score at learning rate 0.01 and max_depth 10 is 0.8954534315038922

Test f1 score at learning rate 0.01 and max_depth 10 is 0.8046964171679296

Train f1 score at learning rate 0.01 and max_depth 15 is 0.9352927741067034

Test f1 score at learning rate 0.01 and max_depth 15 is 0.8039186867555885

Train f1 score at learning rate 0.1 and max_depth 5 is 0.9083537405579349

Test f1 score at learning rate 0.1 and max_depth 5 is 0.8042766810681934

Train f1 score at learning rate 0.1 and max_depth 10 is 0.9576882082144578

Test f1 score at learning rate 0.1 and max_depth 10 is 0.800920784262196

Train f1 score at learning rate 0.1 and max_depth 15 is 0.9591385799663188

Test f1 score at learning rate 0.1 and max_depth 15 is 0.7979809452060369

```
from catboost import CatBoostClassifier
n_estimators = [100,200,500]
class_weights = [{0:1,1:2},{0:1,1:4},{0:1,1:6}]

#Loop for the Tuning
for n in n_estimators:
    for l in class_weights:
        clf = CatBoostClassifier(learning_rate=0.1,max_depth = 15,n_estimators = n,class_weights=class_weights)
        clf.fit(X_train_encode,y_train)

        #y_pred = (clf.predict_proba(X_test)[:,-1] >= 0.21).astype('int')
        y_train_pred = clf.predict(X_train_encode)
        y_test_pred= clf.predict(X_test_encode)
        print(f"Train f1 score at n_estimators {n} and class_weights {l} is {f1_score(y_train_pred,y_train)}")
        print(f"Test f1 score at n_estimators {n} and class_weights {l} is {f1_score(y_test_pred,y_test)}")
        print("-"*50)
```

Train f1 score at n_estimators 100 and class_weights {0: 1, 1: 2} is

0.9261981184750693

Test f1 score at n_estimators 100 and class_weights {0: 1, 1: 2} is 0.7897664765934407

Train f1 score at n_estimators 100 and class_weights {0: 1, 1: 4} is

0.9046092086089003

Test f1 score at n_estimators 100 and class_weights {0: 1, 1: 4} is 0.7760937523544432

Train f1 score at n_estimators 100 and class_weights {0: 1, 1: 6} is 0.885049171172361

```

Test f1 score at n_estimators 100 and class_weights {0: 1, 1: 6} is 0.7682022818589821
-----
Train f1 score at n_estimators 200 and class_weights {0: 1, 1: 2} is
0.9422465020598454
Test f1 score at n_estimators 200 and class_weights {0: 1, 1: 2} is 0.7925053645163538
-----
Train f1 score at n_estimators 200 and class_weights {0: 1, 1: 4} is
0.9290554684186333
Test f1 score at n_estimators 200 and class_weights {0: 1, 1: 4} is 0.7845243611761946
-----
Train f1 score at n_estimators 200 and class_weights {0: 1, 1: 6} is
0.9197642374155665
Test f1 score at n_estimators 200 and class_weights {0: 1, 1: 6} is 0.7789790939795675
-----
Train f1 score at n_estimators 500 and class_weights {0: 1, 1: 2} is
0.9553697729915223
Test f1 score at n_estimators 500 and class_weights {0: 1, 1: 2} is 0.7927979542992917
-----
Train f1 score at n_estimators 500 and class_weights {0: 1, 1: 4} is 0.93748023369878
Test f1 score at n_estimators 500 and class_weights {0: 1, 1: 4} is 0.7846137711812216
-----
Train f1 score at n_estimators 500 and class_weights {0: 1, 1: 6} is
0.9284114589396384
Test f1 score at n_estimators 500 and class_weights {0: 1, 1: 6} is 0.7832205988751294
-----

```

```
from catboost import CatBoostClassifier
```

```
clf = CatBoostClassifier(learning_rate=0.01,max_depth = 15,n_estimators = 100,class_weights={0:1,1:6})
clf.fit(X_train_encode,y_train)
```

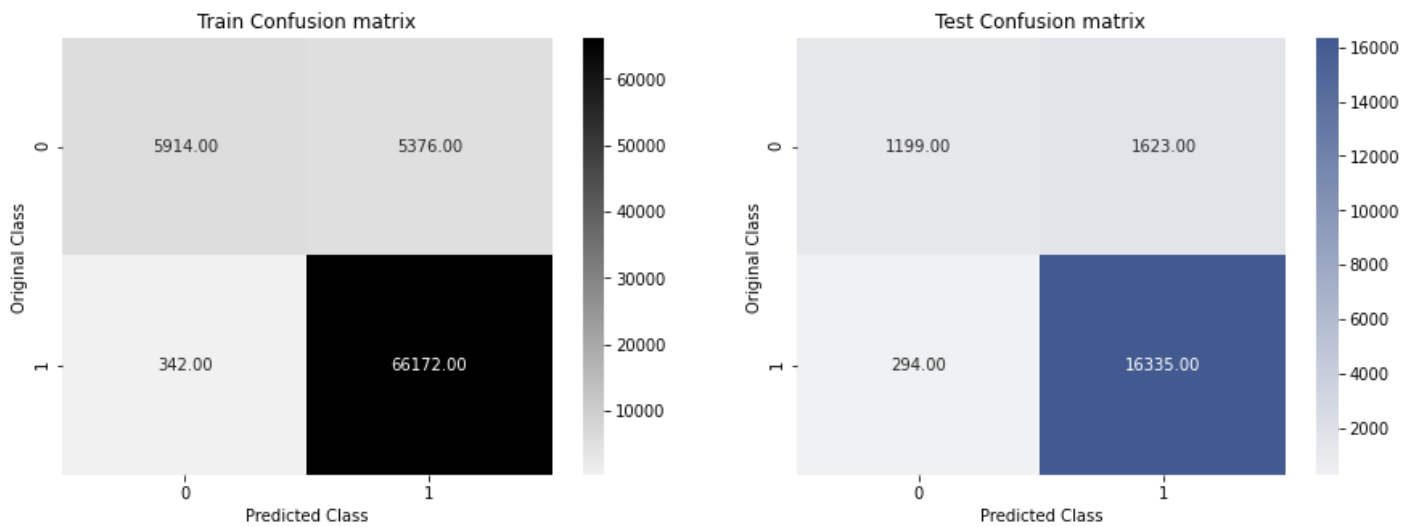
```
<catboost.core.CatBoostClassifier at 0x7f7885bb2f10>
```

```
from sklearn.metrics import f1_score
print('Train f1 score',f1_score(y_train,clf.predict(X_train_encode),average='macro'))
print('Test f1 score',f1_score(y_test,clf.predict(X_test_encode),average='macro'))
```

```
Train f1 score 0.8163487913899979
```

```
Test f1 score 0.7501551782310858
```

```
confusion_matrices_plot(y_train,clf.predict(X_train_encode),y_test,clf.predict(X_test_encode))
```



```
#saving the model
filename = '/content/drive/MyDrive/case study/models/catboost.sav'
joblib.dump(clf, filename)
```

```
['/content/drive/MyDrive/case study/models/catboost.sav']
```

8.StackingClassifier

```
#https://machinelearningmastery.com/stacking-ensemble-machine-learning-with-python/
# defining the base models
est0 = list()
est0.append(('lr', SGDClassifier(loss='log', alpha=0.1, n_jobs=-1, random_state=25)))
est0.append(('svm', SGDClassifier(loss='hinge', alpha=1, n_jobs=-1, random_state=25)))
est0.append(('dt', DecisionTreeClassifier(max_depth = 135, min_samples_leaf = 28, min_s
est0.append(('rf', RandomForestClassifier(max_depth = 13, min_samples_leaf = 49, min_sa
                                     n_jobs=-1)))
est0.append(('xgb', XGBClassifier(n_estimators=123, max_depth=3, subsample= 1, learning_r
                                     n_jobs=-1)))
```

```
est_final= SGDClassifier(loss='log',n_jobs=-1, random_state=25)
```

```
clf = StackingClassifier(estimators=est0, final_estimator = est_final, n_jobs=-1)
```

```
clf.fit(X_train_encode,y_train)
```

```
y_train_pred = clf.predict(X_train_encode)
```

```
y_test_pred = clf.predict(X_test_encode)
```

```
# printing train and test scores
```


```
print('Train f1 score', f1_score(y_train,y_train_pred,average='macro'))
```

```
print('Test f1 score', f1_score(y_test,y_test_pred,average='macro'))
```

Train f1 score 0.8749240548287187

Test f1 score 0.7995396369898682

```
confusion_matrices_plot(y_train,y_train_pred,y_test,y_test_pred, 'stack.png')
```



```
#saving the model
filename = '/content/drive/MyDrive/case study/models/stackcl.sav'
joblib.dump(clf, filename)
```

```
['/content/drive/MyDrive/case study/models/stackcl.sav']
```

9.VotingClassifier

```
#https://machinelearningmastery.com/stacking-ensemble-machine-learning-with-python/
# defining the base models
est0 = list()
est0.append(('lr', SGDClassifier(loss='log', alpha=0.1, n_jobs=-1, random_state=25)))
est0.append(('dt', DecisionTreeClassifier(max_depth = 135, min_samples_leaf = 28, min_s
est0.append(('rf', RandomForestClassifier(max_depth = 13, min_samples_leaf = 49, min_sa
n_jobs=-1)))
est0.append(('xgb', XGBClassifier(n_estimators=123, max_depth=3, subsample= 1, learning_r
n_jobs=-1)))
```


```
#https://www.geeksforgeeks.org/ml-voting-classifier-using-sklearn/
# Voting Classifier with hard voting
vot_hard = VotingClassifier(estimators = est0, voting ='hard')
vot_hard.fit(X_train_encode, y_train)
y_train_pred = vot_hard.predict(X_train_encode)
y_test_pred = vot_hard.predict(X_test_encode)

# printing train and test scores
print('Train f1 score', f1_score(y_train,y_train_pred,average='macro'))
print('Test f1 score', f1_score(y_test,y_test_pred,average='macro'))
```

Train f1 score 0.8788060839390608

Test f1 score 0.806804323724443

```
confusion_matrices_plot(y_train,y_train_pred,y_test,y_test_pred, 'votclf_1.png')
```



#saving the model

```
filename = '/content/drive/MyDrive/case study/models/vot_hard.sav'
joblib.dump(vot_hard, filename)
```

```
['/content/drive/MyDrive/case study/models/vot_hard.sav']
```

Voting Classifier with hard voting

```
vot_soft = VotingClassifier(estimators = est0, voting = 'soft')
vot_soft.fit(X_train_encode, y_train)
y_train_pred = vot_soft.predict(X_train_encode)
y_test_pred = vot_soft.predict(X_test_encode)
```


printing train and test scores

```
print('Train f1 score', f1_score(y_train, y_train_pred, average='macro'))
print('Test f1 score', f1_score(y_test, y_test_pred, average='macro'))
```

Train f1 score 0.8776384471346554

Test f1 score 0.8051520742812928

```
confusion_matrices_plot(y_train, y_train_pred, y_test, y_test_pred, 'votclf_2.png')
```



```
#saving the model
filename = '/content/drive/MyDrive/case study/models/vot_soft.sav'
joblib.dump(vot_soft, filename)
```

```
['/content/drive/MyDrive/case study/models/vot_soft.sav']
```

Summary

```
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Model", "Hyper parameter", "test f1-Score"]
x.add_row(['logistic ', 'alpha=0.1', '0.802699'])
x.add_row(['Regression', ' ', ' '])
x.add_row([' ', ' ', ' '])
x.add_row(['Linear SVM', 'alpha=1', '0.793024'])
x.add_row([' ', ' ', ' '])
x.add_row(['Decision Tree', 'max_depth=135', '0.802874'])
x.add_row([' ', 'min_samples_leaf= 28,min_samples_split=193', ' '])
x.add_row([' ', ' ', ' '])
x.add_row(['Random Forest', 'max_depth =13, min_samples_leaf=49', '0.803262'])
x.add_row([' ', 'min_samples_split =165,n_estimators=108', ' '])
x.add_row([' ', ' ', ' '])
x.add_row(['XGBclassifier', 'n_estimators=109, max_depth=3', '0.781894'])
x.add_row([' ', 'subsample=0.5,learning_rate=0.1,colsample_bytree=1', ' '])
x.add_row([' ', ' ', ' '])
x.add_row(['LGBMClassifier', 'n_estimators=500, max_depth=10', '0.806333'])
x.add_row([' ', 'subsample=0.5,learning_rate=0.01,colsample_bytree=0.1', ' '])
x.add_row([' ', ' ', ' '])
x.add_row(['AdaBoost', 'n_estimators=500, learning_rate=0.05', '0.798562'])
x.add_row([' ', ' ', ' '])
x.add_row(['CATBoost', 'earning_rate=0.01,max_depth = 15', '0.750155'])
x.add_row([' ', 'n_estimators = 100,class_weights ={0: 1, 1: 4}', ' '])
x.add_row([' ', ' ', ' '])
x.add_row(['StackingClassifier', 'estimators=est0, final_estimator = est_final', '0.79953'])
x.add_row([' ', ' ', ' '])
x.add_row(['VotingClassifier', ' ', ' '])
x.add_row(['vot_hard', 'estimators = est0, voting =hard', '0.80680'])
x.add_row(['vot_soft', 'estimators = est0, voting = soft', '0.80515'])
x.add_row([' ', ' ', ' '])
print(x)
```

```
+-----+-----+-----+
-----+
|      Model      |      Hyper parameter      | test f1-
Score |
+-----+-----+-----+
-----+
|      logistic   |      alpha=0.1      |
0.802699 |
```

Regression	
Linear SVM	alpha=1
0.793024	
Decision Tree	max_depth=135
0.802874	
	min_samples_leaf= 28,min_samples_split=193
Random Forest	max_depth =13, min_samples_leaf=49
0.803262	
	min_samples_split =165,n_estimators=108
XGBclassifier	n_estimators=109, max_depth=3
0.781894	
	subsample=0.5,learning_rate=0.1,colsample_bytree=1
LGBMClassifier	n_estimators=500, max_depth=10
0.806333	
	subsample=0.5,learning_rate=0.01,colsample_bytree=0.1
AdaBoost	n_estimators=500, learning_rate=0.05
0.798562	
CATBoost	earning_rate=0.01,max_depth = 15
0.750155	
	n_estimators = 100,class_weights ={0: 1, 1: 4}
StackingClassifier	estimators=est0, final_estimator = est_final

0.799539			
	VotingClassifier		
	vot_hard		estimators = est0, voting =hard
0.80680			
	vot_soft		estimators = est0, voting = soft
0.80515			
+-----+			
-----+			